DETAILED REPORT IS BELOW –

Phase 2 - Ideate

Diverging: Generating the Idea Set

Rather than jumping to the first fix that came to mind, we forced ourselves through the diverge stage properly, using 2 ideation frameworks from class: the 6 Ideation Triggers - frustrations, anomalies, confluence, orthodoxies, extremes, analogies and Assumption Busting. The goal at this stage was quantity and variety, not quality. We judged nothing until convergence, and numbered every idea as it came out so the group could see the full spread at a glance.

Trigger 1 - Frustrations - problems that actively annoy people

1. A live WhatsApp ping every time the job moves to a new stage: received, diagnosed, approved, in repair, ready.
2. A one-tap approve or decline button attached to every extra cost line, sent before the work starts rather than billed after.
3. A digital job card the customer can open on their own phone, showing what the technician is doing right now instead of what was done 3 hours ago.
4. A plain SMS or automated voice-call fallback for customers who do not use WhatsApp or do not own a smartphone.

Trigger 2 - Anomalies - things that do not quite make sense once you notice them

5. Why does the final bill only appear at pickup? A running bill that updates as each part or labour line is added.
6. Why does the advisor call only when there is extra cost, never to say things are still on track? Short "nothing to report, still on schedule" pings between the real updates.
7. Why can a customer track a food delivery rider in real time but not their own car's repair? A short diagnostic photo or video clip attached to the job card at the point it is taken.

Trigger 3 - Confluence - two or more existing trends colliding

8. Smartphone penetration plus UPI adoption: pay-as-you-approve micro-payments per repair stage instead of one lump sum at collection.
9. QR codes plus WhatsApp: scan a QR code at drop-off to open a live tracking thread tied to that specific job.

10. Video calling plus remote-work habits: an optional 2-minute video walkthrough of the vehicle with the technician before approving a major repair.

Trigger 4 - Orthodoxies - questioning what everyone in the category assumes is fixed

11. Ques: "the service advisor is the only channel."

    Ans: a simple chatbot handles routine status questions so the advisor is only pulled in for judgement calls.

12. Ques: "an estimate is given once, at the start."

    Ans: a tiered estimate like definite, likely, possible shown upfront, so a later revision feels expected rather than sprung on the customer.

Trigger 5 - Extremes - edge-case users who reveal needs the average user cannot articulate

13. A customer who has to leave town right after drop-off: a fully remote approval flow with a digital signature.
14. A customer who has been overcharged before and no longer trusts any service centre: an optional independent photo-verification step for major repairs.
15. A first-time car owner who does not understand technical terms: plain-language notes attached to each diagnostic line, written for someone who has never opened a bonnet.

Trigger 6 - Analogies - borrowing structure from a different industry solving a similar problem

16. From food delivery: a named-stage tracker, Received, Diagnosed, Approved, In Repair, Quality Check, Ready, the same idea as an order tracker.
17. From hospital patient portals: one login that stores every past visit, part replaced, and amount paid, so the customer is not starting from zero each time.
18. From courier tracking: a dynamic ready-by estimate that updates itself the moment a delay is known, instead of one static promise made at drop-off.

Assumption Busting

We also ran the Assumption Busting technique on our own category

List the assumptions we make about service communication: the customer has to physically visit or call to know what is happening; the advisor is the one who decides when to share an update; nothing gets communicated unless there is a problem.

- Reverse one: what if the customer never had to ask at all?
- Explore what that opens up:

19. Automatic push updates the customer never requested, instead of the customer having to check in.
20. A shareable tracking link sent by SMS at drop-off, so no app download is needed.
21. A short automatic health report sent even on visits where nothing went wrong, so silence stops being the default.

We did not filter any of these 21 for realism yet. That comes at convergence, and it is well past the 4 or 5 ideas that come from familiar territory alone like an app, a phone call script, a feedback form which is what the exercise was for.

Converging: Choosing What to Prototype

For convergence we used the 4-lens scoring model from class which includes - desirability, feasibility, viability, novelty rather than a straight vote, because a few of the ideas above were popular with the group for reasons that had nothing to do with whether they would actually work in this sector. We shortlisted 5 ideas that came up repeatedly across different triggers and scored each on a 1-5 scale.

Shortlisted idea	**Desirability**	**Feasibility**	**Viability**	**Novelty**	**Total**
Stage-wise WhatsApp status and approval flow (Ideas 1, 2, 16, 19)	5	4	4	3	16
Dynamic running bill visible during the visit (Idea 5)	4	3	3	3	13
Optional video walkthrough before major repairs (Idea 10)	4	2	2	4	12
One login storing full-service history (Idea 17)	3	3	4	2	12
Independent third-party photo verification (Idea 14)	3	2	2	4	11

The stage-wise WhatsApp flow scored highest and, on inspection, is really a container for several of the smaller ideas above: the approve/decline buttons from Idea 2, the named-stage tracker borrowed in Idea 16, and the automatic "nothing to report" pings from Idea 6 and Idea 19.

We are carrying this forward as our concept: a WhatsApp-based thread, opened by QR code at drop-off, that sends the customer a message at every job-stage change and requires an explicit tap to approve any cost beyond the original estimate.

We are carrying the video-walkthrough idea (Idea 10) forward as a secondary, opt-in feature rather than a separate finalist, since it tested well on novelty but depends on technician time we cannot assume every centre has spare.

Revisiting Sector Constraints

This concept works inside 3 of our 4 constraints without a fight. It needs no new hardware, capital intensity stays low, since it runs on WhatsApp Business tooling rather than a bespoke app and it does not touch the OEM's warranty disclosure process, regulatory risk stays low and a wrongly worded status message is an inconvenience rather than a safety issue like failure cost stays manageable, provided diagnostic content is still written and checked by a human technician, not generated automatically.

It does have to work around cycle time. Trust in this category is built over repeat visits, so a single good experience with the new flow will not, by itself, be strong evidence that customers have started trusting the service centre again. That is exactly why **Phase 4** looks at behavior and repeat intent, not just a one-time reaction.

Phase 3 - Prototype

The Riskiest Assumption

The idea itself is not really in question, customers clearly want more visibility. What we do not know is whether they will actually act on it: will a vehicle owner tap "approve" on a WhatsApp message from an unfamiliar number as readily as they would answer a phone call from their usual advisor, or will they ignore the message, call the service centre to ask "is this real," and end up needing the phone call anyway? If they do not act on the messages, the new channel adds a way to communicate without solving the underlying trust problem.

Riskiest assumption: vehicle owners will read, trust, and respond to stage-wise WhatsApp approval requests quickly enough that the process feels faster and clearer than a phone call, not slower.

Choosing the Prototype Level

Using the Prototype Ladder from class, we deliberately stayed low rather than building a working bot. A clickable mockup of the chat thread would show us whether people understand the interface, but it cannot tell us whether a real vehicle owner, mid-service, will actually respond within a useful window, since a mockup carries none of the pressure or distraction of an actual repair day. Building a real WhatsApp Business API integration, on the other hand, would be a level 6 or 7 commitment for an assumption we have not validated yet, which is exactly the mistake the "idea trap" case studies warned against.

We chose a Wizard-of-Oz pilot instead: 2 group members manually sent WhatsApp-style status and approval messages, timed to match the technician's real progress on the vehicle, while a third observed and timed the customer's responses. The customer experiences something that looks and feels automated; behind the scenes, it is entirely manual. This is the cheapest version that can still answer our real question, which is about behaviour, not interface design.

Experiment Card

Field	Detail
Assumption	We believe that vehicle owners will respond to stage-wise WhatsApp approval requests without needing a follow-up phone call, and will trust the process at least as much as a call from their advisor.
Test	A Wizard-of-Oz pilot at a local multi-brand garage and one authorised service centre, run manually over real service visits.
Metric	Time to respond to each approval request; percentage of approvals given without a follow-up call to the centre; a 1-5 trust rating collected at pickup.
Success threshold	At least 70% of approval requests answered within 30 minutes without a follow-up call, and an average trust rating of 4 or higher.
Result	See Phase 4.
Decision	Persevere, iterate, pivot, or kill, based on the result against the threshold above.

We filled in the success threshold before running a single test, for the reason given in class: it stops us from quietly redefining "success" after we have already seen how the pilot went.

Sector Constraint Note

Running even a Wizard-of-Oz version of this test meant asking a real garage owner and one authorised advisor for permission to message their customers, which took longer to arrange than the test itself, a small but real reminder of the cycle-time constraint. We also had to be careful with customer phone numbers collected for the pilot, storing them only for the duration of the test and deleting them afterward, since any future real version of this concept will need to think seriously about data protection before it touches customer WhatsApp numbers at scale.

Phase 4 – Test

Running the Test

We ran the pilot with 8 real vehicle owners across the campus, none of them teammates or people we knew personally. All 8 had booked a routine service or a minor repair in the normal course of their week; none were recruited specifically for this study, which meant we caught them in an ordinary, slightly rushed state of mind rather than a reflective one.

Rather than asking "would you use this", we watched what people actually did with the messages: whether they opened them, how long they took to reply, and whether they picked up the phone anyway. We then used the I Like / I Wish / What If / I Noticed structure at pickup to collect structured feedback instead of a plain opinion.

What We Heard

Category	What came up
I Like	Knowing the cost before it changed, rather than being told at the counter. Several people specifically liked seeing "still on track" messages on visits where nothing had gone wrong.
I Wish	2 customers wished the messages had come from a saved contact name rather than an unknown number, since one almost mistook the first message for spam. One wished for a Hindi-language option.
What If	1 customer suggested the same thread could hold the next service reminder automatically. Another asked whether the thread could show the technician's name and photo, to make the message feel less anonymous.
I Noticed	We noticed that people responded fastest to the approval message with a rupee amount attached, and slowest to purely informational updates, several of those were only opened once the customer arrived to collect the car.

Reading the Evidence

Against the Evidence Ladder from class, our result sits above simple opinion but below proof through repeat behaviour, since a single visit cannot show repeat use. 6 of the 8 customers responded to the cost-approval message within 30 minutes without calling the centre, 1 responded within 40 minutes, and 1 never opened the message and only found out about the extra cost at pickup, the exact failure we set out to fix. The average trust rating collected at pickup was 4.1 out of 5. That clears the response-time and trust thresholds we set in the Experiment Card, though narrowly, and the one customer who missed the message entirely because it arrived from an unsaved number and looked, to them, like spam is a real weak point, not a rounding error.

Decision: Persevere / Pivot / Kill

Our core assumption held: most customers did respond quickly, without a follow-up call, and rated the experience as trustworthy. We are persevering with this concept, with one concrete change carried into Phase 5, the messages need to come from a verified, named business account rather than a personal number, since the near-miss with the spam-looking message is exactly the kind of trust problem the whole concept was meant to solve, not create.

Phase 5 - Scale

Scale Readiness Check

Dimension	Question	Where we stand
Desirability	Are customers repeatedly choosing it?	Strong signal from one visit each; we do not yet have repeat-visit data, since our pilot customers have not been back for a second service.
Feasibility	Can we deliver it reliably?	Yes, at small scale with manual messaging. A real rollout needs a verified WhatsApp Business account and a lightweight internal tool so advisors are not typing every update by hand.
Viability	Do the economics and resources work?	Cost per centre is low (messaging fees plus a small amount of advisor time), but that time saving has to be measured against the extra minute or two advisors spend keying in updates, which we have not yet quantified.
Repeatability	Does it work across users, locations, or contexts?	Untested outside two locations in the same city. Advisor discipline varied even between our two pilot sites, one centre's staff sent updates promptly, the other needed reminding twice during the pilot.

Desirability and feasibility look ready to move past pilot stage; viability and repeatability are the 2 dimensions we would want more evidence on before recommending a network-wide rollout.

Sector Constraints, Revisited Again

Taking this from 2 garages to a real network changes what each constraint demands. Cycle time means we would want at least 2 full service cycles per customer, so 3 to 6 months per pilot centre, before treating trust improvements as real rather than a first-visit courtesy effect. Capital intensity stays favourable since the core solution is still software, though a multi-centre rollout would need a shared WhatsApp Business API subscription and a simple dashboard for advisors, a modest but real cost compared to the near-zero cost of our manual pilot. Regulatory requirements grow the most: storing

customer phone numbers and message logs across multiple locations means we would need an actual data-retention and consent policy before scaling, not just the ad hoc care we took during the pilot. Failure cost stays contained as long as every diagnostic message is still written or approved by a human technician rather than auto-generated, which is a rule we would want written into the rollout plan explicitly, not left as an assumption.

What Stays Constant, What Changes

What stays constant if this expanded to a new city or a different OEM's network: the underlying operating principle, that the customer is told the diagnosis and cost before being asked to pay, not after; and the requirement that every cost change needs an explicit tap to approve, not a phone call after the fact.

What would need to change: the channel itself, since WhatsApp is the obvious default in most Indian cities but would need an SMS or IVR fallback in areas with lower smartphone penetration; the language and tone of message templates, since our pilot ran only in English and one customer specifically asked for Hindi; and the warranty-disclosure wording inside each message, since that has to match whichever OEM's dealership network the service centre belongs to, which is a regulatory constraint we flagged as far back as Phase 1 and have not been able to design around, only design within.

Summary

Across Phases 2 through 5, the group moved from 21 raw ideas to 1 tested concept: a stage-wise WhatsApp status and approval thread built to close the exact communication gap identified in Phase 1, and to directly address our chosen problem, the lack of trust in vehicle servicing. A Wizard-of-Oz pilot with 8 real vehicle owners gave us evidence, not just opinion, that the core assumption holds: most customers will respond quickly and trust the process, provided the messages come from a verified source. Before this could scale past two garages, we would want more repeat-visit data on desirability, a clearer read on advisor time cost for viability, and a proper data-retention policy to satisfy the regulatory constraint that has shaped this problem from the start.

AI Disclosure

Claude was used to help structure and draft this write-up from the group's own ideation notes, prototype observations, and pilot interview records, and to score the shortlisted ideas against the 4-lens criteria as described in the Ideate section. The ideas generated, the prototype design choices, the pilot results, and the final Persevere / Pivot / Kill and scale-readiness judgements reflect the group's own analysis and fieldwork. AI was used for support and alignment, not for generating the underlying evidence or decisions.